Ameen Salim

925-858-5314 | ameensalim1@gmail.com | linkedin.com/in/ameensalim/ | github.com/ameensalim1

Education

Georgia Institute of Technology — Atlanta, GA (Remote)
M.S. in Computer Science — *Expected Graduation: June 2027*

University of California, Davis — Davis, CA
B.S. in Computer Science — *June 2025*

- Relevant Coursework: Data Structures and Algorithms, Operating Systems, Advanced Algorithms, Systems Programming, Linear Algebra, Computer Architecture, Machine Learning, Scientific Computation, Deep Learning

Technical Skills

Languages: Python(3yrs), JavaScript(2yrs), C/C++(1yrs), SQL
Software: GCP, Docker, React.js, Node.js, Flask, FastAPI, React Native, MongoDB, Express.js, Next.js, NLP, pandas, NumPy, ffmpeg, PyTorch, sci-kitlearn, OpenCV, GitHub CI, LLMs, Retrieval-Augmented Generation, Agentic AI, Git

Experience

Software Engineer — June 2024 – Present
MyGenie — *Bay Area, CA*

- Building automated AI training pipeline for RadNerf model avatars, replacing manual processes and reducing time-to-deployment.
- Engineered video processing API to train videos via GPU-accelerated Google Cloud Run.
- Built APIs deployed on GCP using GCS and Cloud Run, enabling seamless video-to-avatar training with a 25% increase in throughput and a 40% reduction in processing time.
- Optimized facial extraction algorithms by tweaking parameters and adding pixel offsets for better orientation detection, improving accuracy by 15% and significantly reducing retries in the training process.

Software Engineer Associate — September 2024 – June 2025
Artificial Intelligence Student Collective — *Davis, CA*

- Architecting and developing WebSocket API backend for real-time ASL translation, achieving sub-100ms latency for end-to-end processing in 5-person engineering team.
- Implementing efficient computer vision pipeline using MediaPipe for hand landmark detection, processing 30+ frames per second with 95% landmark detection accuracy.
- Optimized video frame encoding/decoding pipeline reducing payload size while preserving image quality for ML model inference.

AI Software Engineer Intern — June 2024 – September 2024
PrimisAI — *Los Gatos, CA*

- Debugged and tested a multi-AI agent system for RTL code review and correction, ensuring database retrieval consistency and hardware-specific topic accuracy. Validated retrieval system behavior for FPGA-specific IP catalogs used in 10k+ interactions, improving pass@1 accuracy from 60% to 77% in Verilog generation benchmarks.
- Benchmarked metrics for RAG Pipeline System, increasing system efficiency by 20%, reducing processing time, and enhancing retrieval accuracy.
- Spearheaded CI/CD pipeline implementation using GitHub Actions, accelerating API release cycles and minimizing downtime during updates.

Projects

Analyzing Policing Efficiencies in San Francisco | *Python, TensorFlow, sci-kitLearn*

- Operationalized a machine learning model to optimize police resource allocation by predicting incidents and processing a dataset of 900k+ dispatch records from the San Francisco Police Department, with the goal of reducing response times and enhancing public safety.
- Utilized Python with pandas, scikit-learn, and TensorFlow to implement Random Forest and K-means clustering, achieving an MSE of 59.41 and R-squared of 0.996 for monthly incident predictions.

Loan Approval Predictor | *Python, Pandas, sci-kitLearn*

- Used a 100k+ financial dataset to evaluate loan approval chances based on features such as number of assets, CIBIL score, and credit score.
- Engineered and validated two predictive machine learning models, achieving accuracies of 91.9% and 89.4% through hyperparameter tuning and cross-validation.