Ameen Salim

925-858-5314 | ameensalim1@gmail.com | linkedin.com/in/ameensalim/ | github.com/ameensalim1

Education

Georgia Institute of Technology — Atlanta, GA (Remote)
M.S. in Computer Science — *Expected Graduation: June 2027*

University of California, Davis — Davis, CA
B.S. in Computer Science — *June 2025*

- Relevant Coursework: Machine Learning, Deep Learning, Data Structures & Algorithms, Advanced Algorithms, Operating Systems, Systems Programming, Linear Algebra, Computer Architecture, Scientific Computation

Technical Skills

Languages: Python, C/C++, SQL, TypeScript
Core ML/AI: PyTorch, scikit-learn, Pandas, NumPy, Transformers, HuggingFace, OpenCV
Agentic/GenAI Systems: LLMs, Retrieval-Augmented Generation (RAG), Agentic AI Systems
MLOps/Cloud: Docker, GCP(GCS, Cloud Run, Vertex AI), AWS (S3), CI/CD (GitHub Actions), RESTful APIs (FastAPI)

Experience

AI Software Engineer Intern — June 2024 – September 2024
PrimisAI — *Los Gatos, CA*

- Enhanced a multi-agent system for RTL code generation by improving agentic orchestration and tool utilization, boosting Verilog generation pass@1 accuracy by **17%** (from 60% to 77%).
- Optimized a textual **RAG system** by benchmarking and refining the information retrieval pipeline, reducing LLM processing time and improving overall system efficiency by **20%**.
- Developed a robust CI/CD pipeline using GitHub Actions for automated testing and deployment of ML services, improving developer velocity and reducing release cycle time.

Software Engineer — June 2024 – Present
MyGenie — *Bay Area, CA*

- Designed and deployed a scalable, **production-ready ML system** on GCP to automate a 3D avatar generation pipeline, reducing manual effort and accelerating model deployment cycles.
- Built GPU-accelerated video processing APIs using Docker and Cloud Run, increasing model training throughput by **25%** and cutting data processing time by **40%**.
- Partnered with the **MLOps** team to automate post-training model versioning and artifact storage in cloud databases (Postgres, S3), ensuring data consistency and pipeline integrity.
- Improved ML training pipeline reliability by over **90%** by implementing stateful, context-aware mechanisms that minimize training failures and manual restarts.

Software Engineer Associate — September 2024 – June 2025
Artificial Intelligence Student Collective — *Davis, CA*

- Deployed a low-latency (sub-100ms) WebSocket API using FastAPI for a real-time American Sign Language translation model, enabling seamless multi-turn conversationality.
- Engineered a computer vision data processing pipeline using OpenCV and MediaPipe, achieving **95%** accuracy in hand landmark detection at over **30 FPS** for robust, real-time model inference.
- Optimized video frame encoding to reduce network payload size by **50%** while maintaining image fidelity, ensuring scalable and accurate ML inference for downstream applications.

Projects

Analyzing Policing Efficiencies in San Francisco | *Python, PyTorch, scikit-learn, Pandas*

- Developed and validated a predictive model to forecast crime incidents using over 900k dispatch records, enabling optimized police resource allocation strategies for the SFPD.
- Executed the full ML lifecycle from data exploration with **Pandas** to model experimentation, implementing Random Forest and K-means models in **scikit-learn** to achieve an R-squared of **0.996**.

Loan Approval Predictor | *Python, Pandas, scikit-learn*

- Built a classification model using financial data for over 100k applicants to predict loan approval viability, informing risk assessment processes.
- Achieved **91.9%** classification accuracy by leveraging **scikit-learn** for comprehensive feature engineering, hyperparameter tuning, and cross-validation.